

README

Neural Bending Toolkit

Interface-Level Operations · Logging · Atlas Builder · Semantic Drift Engine

This repository contains a modular toolkit for exploring neural bending through interface-level manipulations, recursive operations, drift generation, and empirical analysis.

It includes:

`interface_bends.py` — Core bend engine (OpenAI + local HuggingFace)

`config.yaml` — Profiles, defaults, local model settings

`process_runs.py` — Log explorer & exporter

`atlas_builder.py` — Appendix-ready Run Atlas generator

`drift_score.py` — Semantic Drift Score calculator

`runs/` — Automatically generated logs for all experiments

The aim is to produce reproducible, research-grade field logs that visualize and quantify interface-level operations described in the Neural Bending Manual.

1. Installation

Create venv (Mac Intel / Linux / Windows)

`python3 -m venv .venv`

`source .venv/bin/activate`

Install dependencies

`pip install --upgrade pip`

`pip install -r requirements.txt`

Required packages

`requirements.txt` should include:

`python-dotenv`

`openai>=1.0.0`

`pyyaml`

`transformers`

`scikit-learn`

For local HuggingFace models

Install PyTorch separately (Mac Intel):

`pip install torch`

2. Configuration (config.yaml)

All runtime behavior is controlled through config.yaml.

Key parts:

```
provider: openai # or: hf_local  
profile: default # default profile used unless overridden
```

```
openai:  
model: gpt-4.1-mini  
max_tokens: 512  
api_key_env: OPENAI_API_KEY
```

```
hf_local:  
model_name: gpt2  
max_new_tokens: 256  
device: cpu  
temperature: 1.0
```

```
profiles:  
default:  
temperature: 1.0  
max_tokens: 512
```

```
clean:  
temperature: 0.4  
entropy_seed_runs: 3  
prompt_recursion_steps: 3  
jitter_min_delay: 0.01  
jitter_max_delay: 0.06
```

```
noisy:  
temperature: 1.3  
entropy_seed_runs: 8  
prompt_recursion_steps: 6  
jitter_min_delay: 0.02  
jitter_max_delay: 0.30
```

Override profile at runtime:

```
--profile clean  
--profile noisy
```

3. Running Interface-Level Bends (interface_bends.py)

The core script performs 5 interface-level bends:

Null Prompt

Prompt Recursion

Entropy Seed

Context Collapse

Interface Jitter

All experiments automatically log results to runs/ as JSON.

Global flags must come before subcommands:

```
python interface_bends.py [GLOBAL FLAGS] SUBCOMMAND [SUBFLAGS]
```

3.1 Null Prompt

Empty input or contradictory instruction.

```
python interface_bends.py --profile noisy --tag test1 null-prompt --mode empty  
python interface_bends.py --tag silence null-prompt --mode contradictory
```

3.2 Prompt Recursion

Outputs become next inputs.

```
python interface_bends.py --profile clean  
--tag motel-case  
prompt-recursion  
--seed-prompt "Describe a missing room in a building."
```

3.3 Entropy Seed

Multiple generations with temperature/seed variance.

```
python interface_bends.py --profile noisy  
--tag raf-seq  
entropy-seed  
--prompt "Write a short paragraph about neural bending."
```

3.4 Context Collapse

Remove or scramble hierarchy.

```
python interface_bends.py --tag collapse1  
context-collapse  
--prompt "Explain what you are."
```

With metadata scramble:

```
--scramble-meta --drop-system-in-scramble
```

3.5 Interface Jitter

Simulated machine-speech pacing.

```
python interface_bends.py --profile clean  
--tag jitter-night  
interface-jitter  
--prompt "Describe a city through security cameras."
```

4. Run Log Structure

Every experiment generates:

```
runs/YYYYMMDDThhmssZ_command_tag.json
```

Contents:

```
{  
  "timestamp": "...",  
  "command": "entropy-seed",  
  "meta": {  
    "tag": "raf-seq",  
    "provider": "openai",  
    "model": "gpt-4.1-mini",  
    "max_tokens": 512  
  },  
  "iterations": [  
    {  
      "run": 1,  
      "prompt": "...",  
      "output": "...",  
      "temperature": 1.18,  
      "seed": 42  
    }  
  ]  
}
```

5. Log Post-Processor (process_runs.py)

Explore, export, and reorganize run logs.

5.1 List runs

```
python process_runs.py list  
python process_runs.py list --tag raf-seq  
python process_runs.py list --command-filter entropy-seed
```

5.2 Export runs to Markdown

```
python process_runs.py export-md  
--tag raf-seq  
--command-filter entropy-seed  
--out exports/raf-seq.md  
--include-prompts  
--include-outputs
```

5.3 Export runs to CSV

```
python process_runs.py export-csv  
--out exports/all-runs.csv
```

5.4 Generate a diagram prompt

For your diagram-generation model:

```
python process_runs.py diagram-prompt  
--tag raf-seq  
--command-filter entropy-seed
```

6. Appendix Builder (atlas_builder.py)

Generates a polished Run Atlas suitable for a dissertation appendix.

Build an atlas for all Interface-Level bends:

```
python atlas_builder.py  
--out appendix/interface-level-atlas.md  
--include-diagram-prompt
```

Build an atlas for a specific tag:

```
python atlas_builder.py  
--tag raf-seq  
--out appendix/raf-seq-atlas.md
```

Build an atlas for a specific bend type:

```
python atlas_builder.py  
--command-filter entropy-seed  
--out appendix/entropy-atlas.md
```

7. Semantic Drift Score (drift_score.py)

Computes semantic drift metrics over logged outputs, using TF-IDF cosine similarity.

Drift is defined as:

```
drift = 1 - cosine_similarity
```

Compute drift for all runs:

```
python drift_score.py
```

Filter by tag:

```
python drift_score.py --tag raf-seq
```

Filter by bend type:

```
python drift_score.py --command-filter entropy-seed
```

Export drift table to CSV:

```
python drift_score.py  
--tag raf-seq  
--command-filter entropy-seed  
--out-csv exports/raf-seq-drift.csv
```

Produces:

```
==== Semantic Drift Summary (per run) ===
```

idx | timestamp | cmd | tag | n_it | pair_drift | seq_drift

0 | 20251122T1245Z | entropy-seed | raf-seq | 5 | 0.372 | 0.428

==== Global Drift ===

Global mean drift: 0.358

8. Project Structure

```
.  
|   ├── interface_bends.py  
|   ├── process_runs.py  
|   ├── atlas_builder.py  
|   ├── drift_score.py  
|   ├── config.yaml  
|   ├── requirements.txt  
|   ├── runs/  
|   |   └── ... auto-generated logs ...  
|   ├── appendix/  
|   ├── exports/  
|   └── README.md
```

9. Summary

This toolkit establishes a fully instrumented foundation for:

empirical neural bending

reproducible interface-level experiments

systematic drift and deformation analysis

diagrammatic documentation

dissertation-appendix production workflows

It formalizes the Neural Bending Manual into a programmable practice.